

Advancements in generative media with Google DeepMind and Google Cloud

Google
Cloud
Next 26





**Alisa
Fortin**

Product Manager,
Google Cloud



**Guillaume
Vernade**


Developer Advocate,
Google Cloud



**Matt
Koenig**

Senior Staff Product
Manager - AI Agents,
Shopify

Gemini inputs and outputs




*Open your phone. Any website. Any
textbook.*

*Media is how we entertain, communicate,
and learn.*

That's what we're building.

genMedia Models



Audio models



Veo



Veo Models: Imagining, prototyping, and assisting at scale

Pricing		Veo 3.1 Lite Cost efficient video with audio	Veo 3.1 Fast Fast video with audio	Veo 3.1 Standard video with audio
720p	per second in USD (\$)	\$0.05	\$0.10	\$0.40
1080p	per second in USD (\$)	\$0.08	\$0.12	\$0.40
4k	per second in USD (\$)	—	\$0.30	\$0.60

Veo 3.1 Lite:

Imagining, prototyping, and assisting at scale



Tips for building with Veo

- ❶ Write descriptive prompts with characters, action, context, camera movement, lighting, atmosphere, emotional mood, etc... (ALS app)
- ❷ Capitalize on advanced reference features
- ❸ Use JSON prompting

Veo Prompting tips

- ① Vividly describe subject and action
- ① Dictate composition and camera motion
- ① Set the lighting and ambiance
- ① Write explicit prompts for audio

Nano- Banana



Nano Banana 2

Faster & Cheaper

5s faster, 50% lower cost. 512px option.

Vibrant & Punchy

High-contrast photography that pops.

Traditional Realism

Crisp textures, efficient processing.

Digital Graphic Assets

Isolated assets with high text accuracy.

Nano Banana Pro

Premium Quality

Optimized for top-tier image quality.

Natural & Muted

Earthy tones and natural aesthetics.

Photorealistic Mastery

Deep shadows and authentic textures.

Physical Design Context

Mimics physical prints and UI screens.

Aspect ratios and 521px to 4K outputs for any project



4k

2k

1k


512px



Proprietary

Search Grounding

Search the web then generate an image of isometric perspective, detailed pixel art that shows the career of Guillaume Vernade



Grounding with Google Search for images

Prompt: generate a watercolor of the Ayub bridge
What? A bridge in Sukkur, Pakistan

Actual bridge



Nano-Banana Pro
(Search grounding only)



Nano-Banana 2
(Search + Image Grounding)

Maps Imagery Grounding


Apply for access here!



Tips for building with Nano-Banana

- ① Use multi-turn editing
- ② Iterate on low resolution then upscale
- ③ Use references
- ④ Get transparent backgrounds (ALS app)
- ⑤ Generate multiple images at once (ALS app)

Lyria 3



Lyria 3 and Lyria 3 Pro

Track Length

Lyria 3 generates 30-second long songs, ideal for rapid prototyping and short-form assets. Lyria 3 Pro supports compositions up to three minutes long.

Vocal Support

Supports vocal conditioning in eight languages (English, German, Spanish, French, Hindi, Japanese, Korean, and Portuguese).


Controls and Conditioning

Lyria 3 Pro includes advanced controls for timed lyrics and tempo control through natural language descriptions.

Multi-modal inputs

You can generate music using text, PDF files, or up to 10 reference images.

Lyria 3 Lite/Pro



Made in
Lyria Studio

Lyria RealTime



SpaceDJ demo

Lyria Core Prompting Framework

[Genre & Style] + [Mood & Energy] + [Instrumentation] + [Tempo & Rhythm] + [Vocal Style & Language]
+ [Timed Song Structure] + [Lyrics]

THE PROMPT

"A 1980s-style synth-pop track, tense and suspenseful, featuring shimmering synthesizers and a driving electronic pulse at 120 BPM. A commanding male baritone with a gravelly, soulful voice."

TIMED SONG STRUCTURE

[Intro] (0:00 - 0:15):

Start with a quiet piano and soft synth pads.

[Verse 1] (0:15 - 0:45):

Bring in a clean, driving beat.

[Chorus] (0:45 - 1:15):

Explode into a massive, anthemic wall of sound.

Lyrics:

[Insert lyrics here or describe the theme]

Lyria 3 Lite/Pro



Made in
Lyria Studio

Lyria RealTime



SpaceDJ demo

Live Jukebox

Text-To-Speech



Gemini 3.1 Flash Text-to-Speech

The latest text-to-speech model from Google DeepMind, offering unprecedented controllability, expressiveness, and audio quality.

Granular Audio Tags

Allows natural language commands directly in text to guide vocal style, pace, and delivery.

"Director's Chair" Controls

Scene Direction: Set physical and emotional environments for multi-turn natural interactions.

Speaker Specificity

Cast unique profiles and use inline tags to change expression mid-sentence.

Global Scale

Supports 70+ languages with style and pacing optimizations for 24 core languages.

Tips for building with TTS

- ❶ Follow the right format: Treat your prompt like a script with director's notes (Tone, Speed, Character).
- ❷ Do not use the raw voices, prompt them (AI Studio)
- ❸ Mix prompts to create more than 2 voices (ALS App)

Mixing models together



goo.gle/cookbook-illustration



Shopify



Matt Koenig

Senior Staff Product
Manager - AI Agents



High-quality brand storytelling at scale

Image gen in the merchant's workflow

- Product fidelity first
- Workflow-native: generation happens inside chat, inline editing, and editors
- Harness control: curated prompts, guardrails, and tuning keep every output on-brand

High-quality brand storytelling at scale

- Every surface
- Nano Banana + Shopify's merchant context = high quality results for every budget
- Chat, inline, editor, brand assets: one harness, built into the tools merchants already use

What do you want?



Tell us what you need

- **Veo**
 - No audio option
 - Synthetic voice selection and controls
 - Draft mode
 - Upscale video
 - Multi-turn
- **TTS/Live**
 - Voice replication
- **Lyria**
 - Video-to-music
- **Nano Banana**
 - Music-to-image
 - Transparent backgrounds
 - Mask-based editing/layering





goo.gle/model-FR

Your feedback is greatly appreciated!



Complete the session
survey in the mobile app



Thank you

Google
 Cloud
Next 26

